

Project on Computer Vision using transfer learning

**By,
Sarikaa S,
2018103058,
12-05-2021.**

Problem Statement:

To classify images using convnet and pre-trained VGG16 models.

Modules for Convnets:

- Loading the dataset.
- Creating Directories for training,testing and validation data.
- Defining the model
- Feed the data to the network ,compile and fit the model
- Test the performance
- Plot Training and Validation graphs

Convnets:

Convnet is a class of deep neural network. It is composed of multiple layers of artificial neurons.

Instantiating a small convnet:

```
In [1]: from keras import layers
from keras import models
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```



```
In [2]: model.summary()
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928
Total params: 55,744		
Trainable params: 55,744		
Non-trainable params: 0		

The summary() function displays the layers ,shape and number of parameters of the convnet model.

Adding a classifier on top of the convnet:

```
In [3]: model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

In [4]: model.summary()

Model: "sequential"
Layer (type)          Output Shape         Param #
===== =====
conv2d (Conv2D)        (None, 26, 26, 32)      320
max_pooling2d (MaxPooling2D) (None, 13, 13, 32)    0
conv2d_1 (Conv2D)       (None, 11, 11, 64)     18496
max_pooling2d_1 (MaxPooling2D) (None, 5, 5, 64)    0
conv2d_2 (Conv2D)       (None, 3, 3, 64)      36928
flatten (Flatten)      (None, 576)           0
dense (Dense)          (None, 64)            36928
dense_1 (Dense)        (None, 10)             650
=====
Total params: 93,322
Trainable params: 93,322
Non-trainable params: 0
```

Since the output has 10 labels in the MNIST dataset ,there are 10 neurons in the last Dense layer.

Training the convnet on MNIST images:

```
In [5]: from keras.datasets import mnist
from keras.utils import to_categorical
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11493376/11490434 [=====] - 2s 0us/step

In [6]: train_images = train_images.reshape((60000, 28, 28, 1))
train_images = train_images.astype('float32') / 255
test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(train_images, train_labels, epochs=5, batch_size=64)

Epoch 1/5
938/938 [=====] - 39s 39ms/step - loss: 0.3952 - accuracy: 0.8736
Epoch 2/5
938/938 [=====] - 37s 39ms/step - loss: 0.0524 - accuracy: 0.9832
Epoch 3/5
938/938 [=====] - 37s 39ms/step - loss: 0.0328 - accuracy: 0.9900
Epoch 4/5
938/938 [=====] - 37s 40ms/step - loss: 0.0231 - accuracy: 0.9929
Epoch 5/5
938/938 [=====] - 37s 40ms/step - loss: 0.0173 - accuracy: 0.9945
Out[6]: <tensorflow.python.keras.callbacks.History at 0x24db9924a90>
```

Evaluating the model on the test data:

```
In [7]: test_loss, test_acc = model.evaluate(test_images, test_labels)
test_acc

313/313 [=====] - 3s 8ms/step - loss: 0.0286 - accuracy: 0.9922
Out[7]: 0.9922000169754028
```

The accuracy of the test data is 99.22%.

Model without max-pooling layer:

```
In [8]: model_no_max_pool = models.Sequential()
model_no_max_pool.add(layers.Conv2D(32, (3, 3), activation='relu',
input_shape=(28, 28, 1)))
model_no_max_pool.add(layers.Conv2D(64, (3, 3), activation='relu'))
model_no_max_pool.add(layers.Conv2D(64, (3, 3), activation='relu'))
```



```
In [9]: model_no_max_pool.summary()

Model: "sequential_1"
-----

| Layer (type)      | Output Shape       | Param # |
|-------------------|--------------------|---------|
| conv2d_3 (Conv2D) | (None, 26, 26, 32) | 320     |
| conv2d_4 (Conv2D) | (None, 24, 24, 64) | 18496   |
| conv2d_5 (Conv2D) | (None, 22, 22, 64) | 36928   |


-----  

Total params: 55,744  

Trainable params: 55,744  

Non-trainable params: 0
```

The disadvantages of not using max pooling that :

1. It doesn't learn the spatial hierarchy of features.
2. It has a large number of parameters which will lead to overfitting.

Training a convnet on cat and dog dataset

The cat and dog dataset is downloaded from Kaggle.

Copying images to training, validation, and test directories:

```
In [2]:  
import os, shutil  
original_dataset_dir = 'C:/Users/ACER/Downloads/train'  
base_dir = 'C:/Users/ACER/Downloads/cats_and_dogs_small'  
os.mkdir(base_dir)  
train_dir = os.path.join(base_dir, 'train')  
os.mkdir(train_dir)  
validation_dir = os.path.join(base_dir, 'validation')  
os.mkdir(validation_dir)  
test_dir = os.path.join(base_dir, 'test')  
os.mkdir(test_dir)  
train_cats_dir = os.path.join(train_dir, 'cats')  
os.mkdir(train_cats_dir)  
train_dogs_dir = os.path.join(train_dir, 'dogs')  
os.mkdir(train_dogs_dir)  
validation_cats_dir = os.path.join(validation_dir, 'cats')  
os.mkdir(validation_cats_dir)  
validation_dogs_dir = os.path.join(validation_dir, 'dogs')  
os.mkdir(validation_dogs_dir)  
test_cats_dir = os.path.join(test_dir, 'cats')  
os.mkdir(test_cats_dir)  
test_dogs_dir = os.path.join(test_dir, 'dogs')  
os.mkdir(test_dogs_dir)
```

```
In [4]:  
fnames = ['cat.{}.jpg'.format(i) for i in range(1000)]  
for fname in fnames:  
    src = os.path.join(original_dataset_dir, fname)  
    dst = os.path.join(train_cats_dir, fname)  
    shutil.copyfile(src, dst)
```

```
In [5]:  
fnames = ['cat.{}.jpg'.format(i) for i in range(1000, 1500)]  
for fname in fnames:  
    src = os.path.join(original_dataset_dir, fname)  
    dst = os.path.join(validation_cats_dir, fname)  
    shutil.copyfile(src, dst)
```

```
In [6]:  
fnames = ['cat.{}.jpg'.format(i) for i in range(1500, 2000)]  
for fname in fnames:  
    src = os.path.join(original_dataset_dir, fname)  
    dst = os.path.join(test_cats_dir, fname)  
    shutil.copyfile(src, dst)
```

```
In [8]:  
fnames = ['dog.{}.jpg'.format(i) for i in range(1000)]  
for fname in fnames:  
    src = os.path.join(original_dataset_dir, fname)  
    dst = os.path.join(train_dogs_dir, fname)  
    shutil.copyfile(src, dst)
```

```
In [9]:  
fnames = ['dog.{}.jpg'.format(i) for i in range(1000, 1500)]  
for fname in fnames:  
    src = os.path.join(original_dataset_dir, fname)  
    dst = os.path.join(validation_dogs_dir, fname)  
    shutil.copyfile(src, dst)
```

```
In [10]: fnames = ['dog.{}.jpg'.format(i) for i in range(1500, 2000)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(test_dogs_dir, fname)
    shutil.copyfile(src, dst)

In [11]: print('total training cat images:', len(os.listdir(train_cats_dir)))
print('total training dog images:', len(os.listdir(train_dogs_dir)))
print('total validation cat images:', len(os.listdir(validation_cats_dir)))
print('total validation dog images:', len(os.listdir(validation_dogs_dir)))
print('total test cat images:', len(os.listdir(test_cats_dir)))
print('total test dog images:', len(os.listdir(test_dogs_dir)))

total training cat images: 1000
total training dog images: 1000
total validation cat images: 500
total validation dog images: 500
total test cat images: 500
total test dog images: 500
```

There are 1000 cat and 1000 dog images in the training directory. And there are 500 cat and 500 dog in validation and testing directory

Instantiating a small convnet for dogs vs. cats classification:

```
In [12]: from keras import layers
from keras import models
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

```
In [13]: model.summary()

Model: "sequential"
-----
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_1 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_2 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_3 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten (Flatten)	(None, 6272)	0
dense (Dense)	(None, 512)	3211776
dense_1 (Dense)	(None, 1)	513

```
=====
```

```
Total params: 3,453,121
Trainable params: 3,453,121
Non-trainable params: 0
```

Configuring the model for training:

```
In [14]: from keras import optimizers
model.compile(loss='binary_crossentropy',
optimizer=optimizers.RMSprop(lr=1e-4),
metrics=['acc'])
```

Using ImageDataGenerator to read images from directories:

```
In [16]: from keras.preprocessing.image import ImageDataGenerator
train_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150),
    batch_size=20,
    class_mode='binary')
validation_generator = test_datagen.flow_from_directory(
    validation_dir, target_size=(150, 150),
    batch_size=20,
    class_mode='binary')
```

```
Found 2000 images belonging to 2 classes.
Found 1000 images belonging to 2 classes.
```

```
In [18]: for data_batch, labels_batch in train_generator:
    print('data batch shape:', data_batch.shape)
    print('labels batch shape:', labels_batch.shape)
    break

data batch shape: (20, 150, 150, 3)
labels batch shape: (20,)
```

Fitting the model using a batch generator:

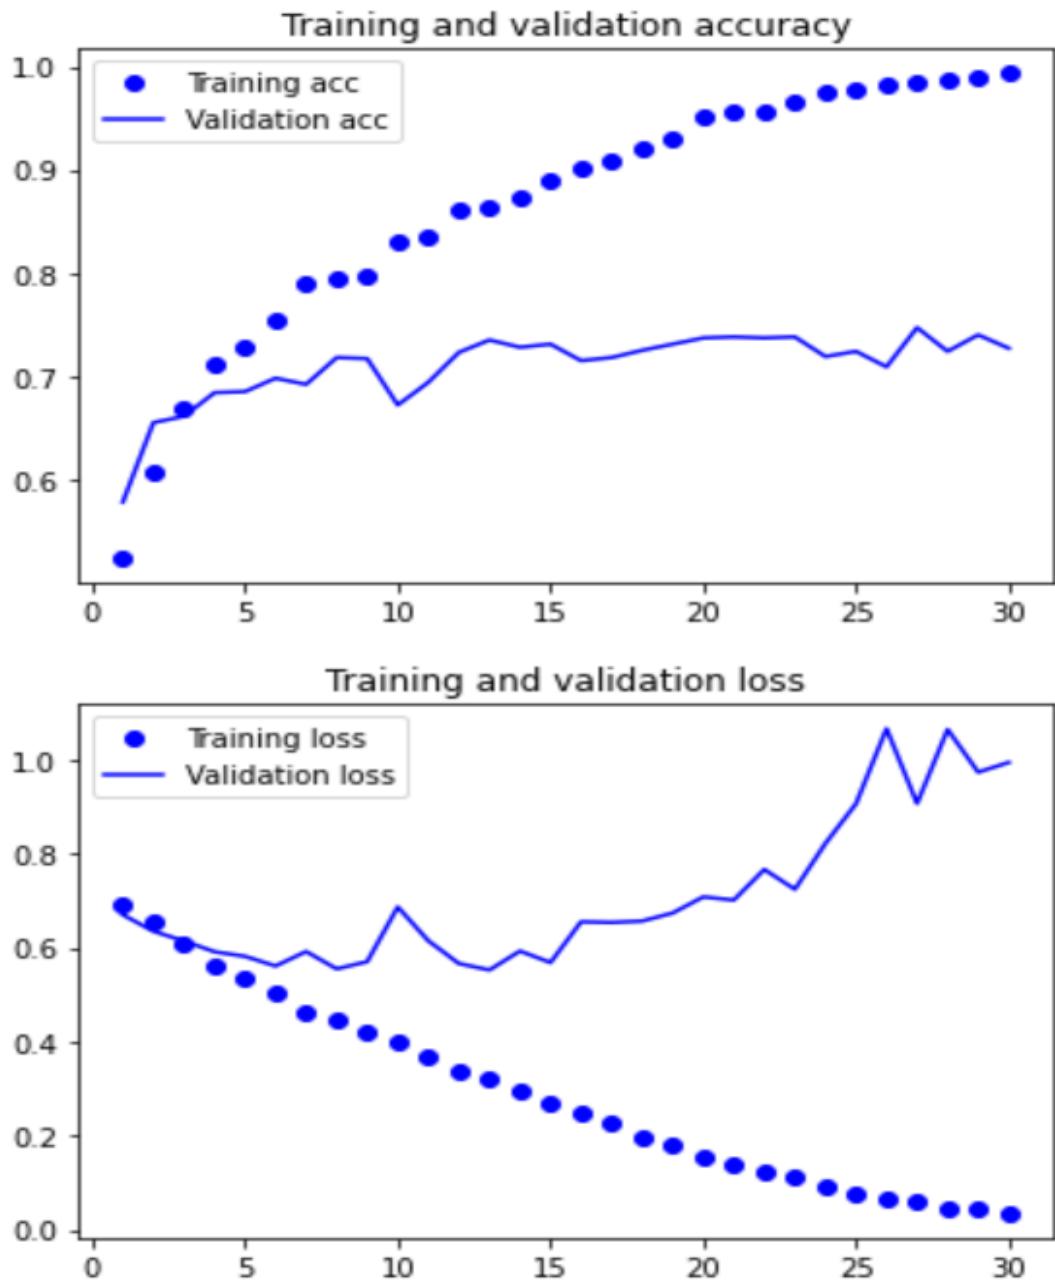
```
In [19]: history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=30,
    validation_data=validation_generator,
    validation_steps=50)

C:\Users\ACER\anaconda3\lib\site-packages\tensorflow\python\keras\engine\training.py:1844: UserWarning: `Model.fit_generator` is
d in a future version. Please use `Model.fit`, which supports generators.
  warnings.warn(`Model.fit_generator` is deprecated and '
```

Epoch 1/30
100/100 [=====] - 59s 576ms/step - loss: 0.7046 - acc: 0.4783 - val_loss: 0.6705 - val_acc: 0.5780
Epoch 2/30
100/100 [=====] - 53s 528ms/step - loss: 0.6647 - acc: 0.5900 - val_loss: 0.6354 - val_acc: 0.6550
Epoch 3/30
100/100 [=====] - 50s 499ms/step - loss: 0.6236 - acc: 0.6656 - val_loss: 0.6147 - val_acc: 0.6610
Epoch 4/30
100/100 [=====] - 50s 501ms/step - loss: 0.5724 - acc: 0.7030 - val_loss: 0.5917 - val_acc: 0.6840
Epoch 5/30
100/100 [=====] - 49s 494ms/step - loss: 0.5348 - acc: 0.7290 - val_loss: 0.5819 - val_acc: 0.6850
Epoch 6/30
100/100 [=====] - 49s 495ms/step - loss: 0.5039 - acc: 0.7664 - val_loss: 0.5613 - val_acc: 0.6980
Epoch 7/30
100/100 [=====] - 49s 492ms/step - loss: 0.4672 - acc: 0.7858 - val_loss: 0.5922 - val_acc: 0.6920
Epoch 8/30
100/100 [=====] - 50s 496ms/step - loss: 0.4462 - acc: 0.8050 - val_loss: 0.5552 - val_acc: 0.7180
Epoch 9/30
100/100 [=====] - 48s 479ms/step - loss: 0.4298 - acc: 0.7845 - val_loss: 0.5706 - val_acc: 0.7170
Epoch 10/30
100/100 [=====] - 52s 518ms/step - loss: 0.3985 - acc: 0.8355 - val_loss: 0.6872 - val_acc: 0.6720
Epoch 11/30
100/100 [=====] - 51s 512ms/step - loss: 0.3734 - acc: 0.8289 - val_loss: 0.6151 - val_acc: 0.6940
Epoch 12/30
100/100 [=====] - 53s 526ms/step - loss: 0.3341 - acc: 0.8502 - val_loss: 0.5667 - val_acc: 0.7230
Epoch 13/30
100/100 [=====] - 50s 494ms/step - loss: 0.3089 - acc: 0.8743 - val_loss: 0.5528 - val_acc: 0.7350
Epoch 14/30
100/100 [=====] - 50s 499ms/step - loss: 0.2680 - acc: 0.8901 - val_loss: 0.5933 - val_acc: 0.7280

Displaying curves of loss and accuracy during training:

```
In [21]: import matplotlib.pyplot as plt
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



Since the gap between training loss curve and validation loss curve is large so overfitting occurs.

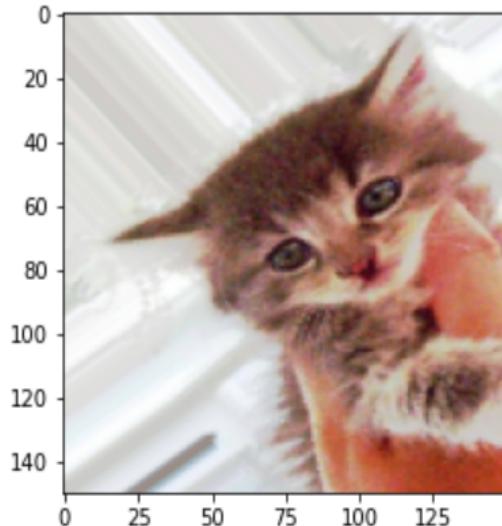
Setting up a data augmentation configuration via `ImageDataGenerator`:

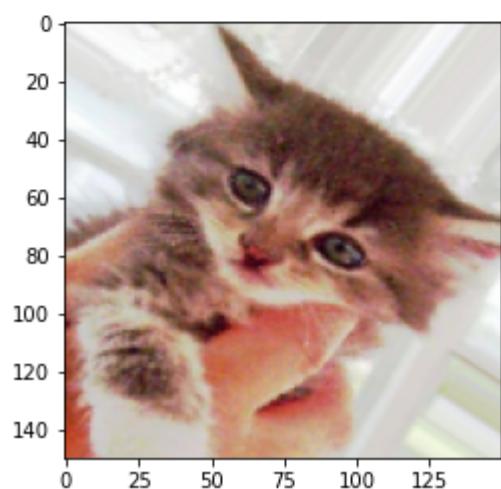
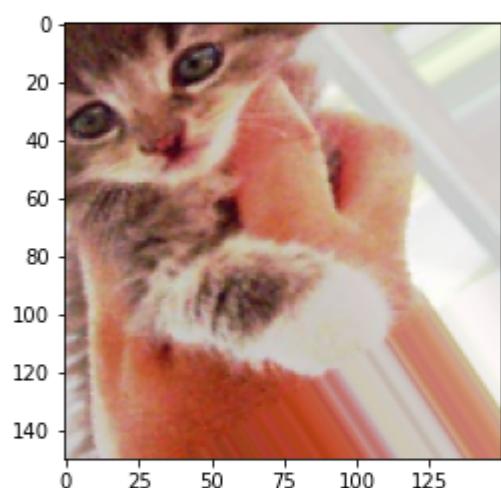
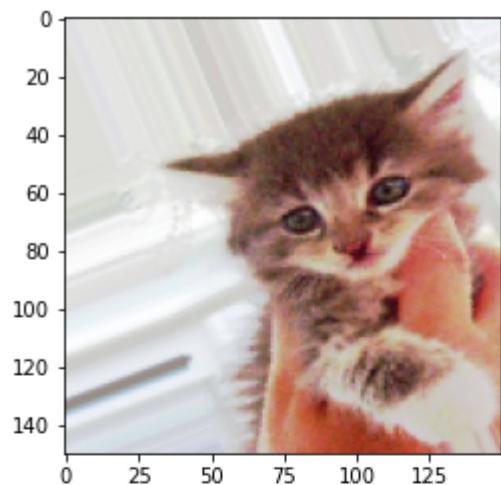
Data augmentation takes the approach of generating more training data from existing training samples, by augmenting the samples via a number of random transformations

```
In [22]: datagen = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest')
```

Displaying some randomly augmented training images:

```
← → ⌂ ⓘ File | C:/Users/ACER/Downloads/Cat%20and%20dog.html  
fill_mode='nearest')  
  
In [23]:  
from keras.preprocessing import image  
fnames = [os.path.join(train_cats_dir, fname) for  
fname in os.listdir(train_cats_dir)]  
img_path = fnames[3]  
img = image.load_img(img_path, target_size=(150, 150))  
x = image.img_to_array(img)  
x = x.reshape((1,) + x.shape)  
i=0  
for batch in datagen.flow(x, batch_size=1):  
    plt.figure(i)  
    imgplot = plt.imshow(image.array_to_img(batch[0]))  
    i += 1  
    if i % 4 == 0:  
        break  
plt.show()
```





Defining a new convnet that includes dropout:

```
In [24]: model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy',
optimizer=optimizers.RMSprop(lr=1e-4),
metrics=['acc'])
```

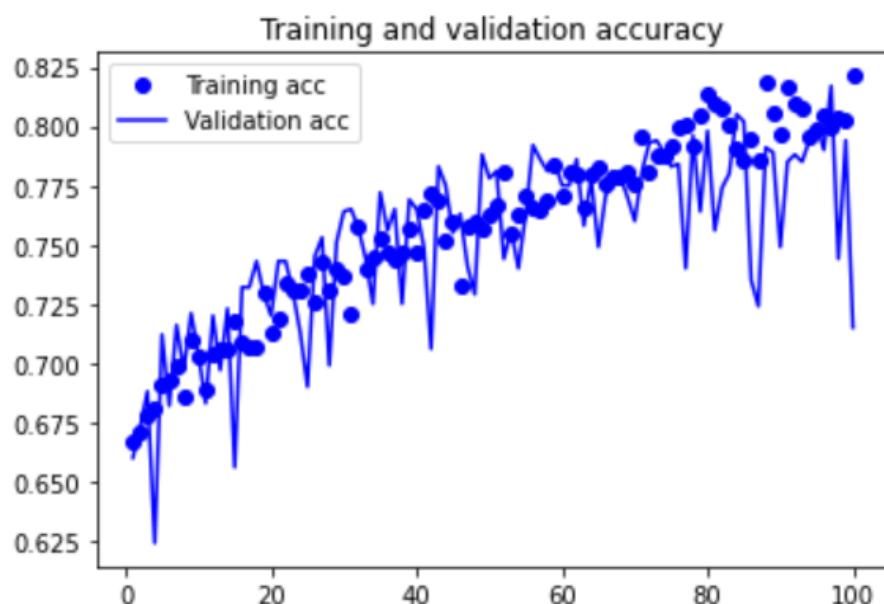
Training the convnet using data-augmentation generators:

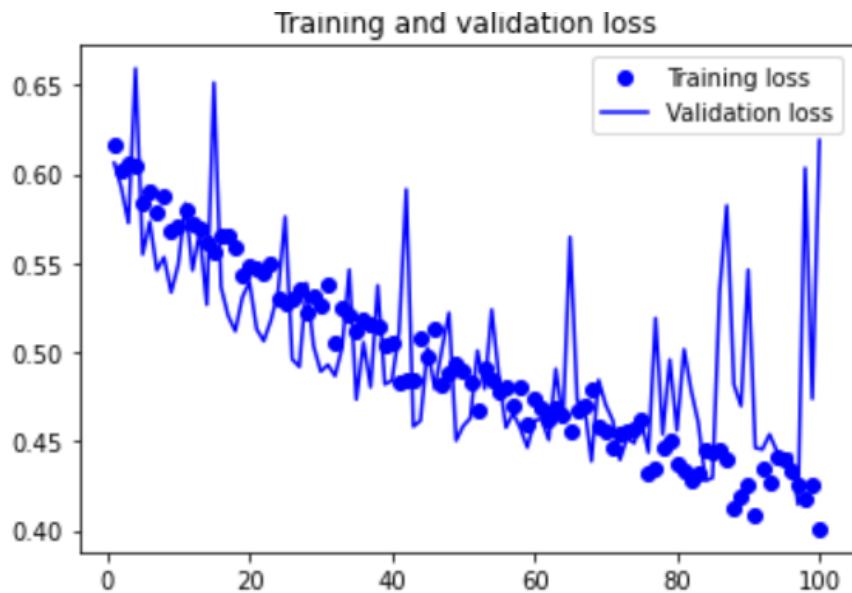
```
In [28]: train_datagen = ImageDataGenerator(
rescale=1./255,
rotation_range=40,
width_shift_range=0.2,
height_shift_range=0.2,
shear_range=0.2,
zoom_range=0.2,
horizontal_flip=True,)
test_datagen = ImageDataGenerator(rescale=1./255)
train_generator = train_datagen.flow_from_directory(
train_dir,
target_size=(150, 150),
batch_size=32,
class_mode='binary')
validation_generator = test_datagen.flow_from_directory(
validation_dir,
target_size=(150, 150),
batch_size=32,
class_mode='binary')
```

Found 2000 images belonging to 2 classes.
Found 1000 images belonging to 2 classes.

```
In [32]: history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=50,  
    epochs=100,  
    validation_data=validation_generator,  
    validation_steps=50)  
  
Epoch 1/100  
50/50 [=====] - ETA: 0s - loss: 0.6160 - acc: 0.6673  
  
Epoch 100/100  
50/50 [=====] - ETA: 0s - loss: 0.4008 - acc: 0.8213
```

```
In [39]: import matplotlib.pyplot as plt  
acc = history.history['acc']  
val_acc = history.history['val_acc']  
loss = history.history['loss']  
val_loss = history.history['val_loss']  
epochs = range(1, len(acc) + 1)  
plt.plot(epochs, acc, 'bo', label='Training acc')  
plt.plot(epochs, val_acc, 'b', label='Validation acc')  
plt.title('Training and validation accuracy')  
plt.legend()  
plt.figure()  
plt.plot(epochs, loss, 'bo', label='Training loss')  
plt.plot(epochs, val_loss, 'b', label='Validation loss')  
plt.title('Training and validation loss')  
plt.legend()  
plt.show()
```





Using a pre trained convnet(Transfer Learning)

Modules:

- Loading the dataset.
- Loading the VGG16 model
- Feed the data to the network ,compile and fit the model
- Test the performance
- Plot Training and Validation graphs

Instantiating the VGG16 convolutional base:

In [34]:

```
from keras.applications import VGG16
conv_base = VGG16(weights='imagenet',
include_top=False,
input_shape=(150, 150, 3))
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58892288/58889256 [=====] - 9s 0us/step

```

← → ⌂ ⌂ File | C:/Users/ACER/Downloads/Cat%20and%20dog.html
In [55]: conv_base.summary()

Model: "vgg16"
Layer (type)          Output Shape         Param #
=================================================================
input_1 (InputLayer)   [(None, 150, 150, 3)]  0
block1_conv1 (Conv2D)  (None, 150, 150, 64)   1792
block1_conv2 (Conv2D)  (None, 150, 150, 64)   36928
block1_pool (MaxPooling2D) (None, 75, 75, 64)  0
block2_conv1 (Conv2D)  (None, 75, 75, 128)   73856
block2_conv2 (Conv2D)  (None, 75, 75, 128)   147584
block2_pool (MaxPooling2D) (None, 37, 37, 128) 0
block3_conv1 (Conv2D)  (None, 37, 37, 256)   295168
block3_conv2 (Conv2D)  (None, 37, 37, 256)   590080
block3_conv3 (Conv2D)  (None, 37, 37, 256)   590080
block3_pool (MaxPooling2D) (None, 18, 18, 256) 0
block4_conv1 (Conv2D)  (None, 18, 18, 512)   1180160
block4_conv2 (Conv2D)  (None, 18, 18, 512)   2359808
block4_conv3 (Conv2D)  (None, 18, 18, 512)   2359808
block4_pool (MaxPooling2D) (None, 9, 9, 512)  0
block5_conv1 (Conv2D)  (None, 9, 9, 512)   2359808
block5_conv2 (Conv2D)  (None, 9, 9, 512)   2359808
block5_conv3 (Conv2D)  (None, 9, 9, 512)   2359808
block5_pool (MaxPooling2D) (None, 4, 4, 512)  0
=====
Total params: 14,714,688
Trainable params: 14,714,688
Non-trainable params: 0

```

Extracting features using the pretrained convolutional base:

```
In [40]:
```

```
import os
import numpy as np
from keras.preprocessing.image import ImageDataGenerator
base_dir = 'C:/Users/ACER/Downloads/cats_and_dogs_small'
train_dir = os.path.join(base_dir, 'train')
validation_dir = os.path.join(base_dir, 'validation')
test_dir = os.path.join(base_dir, 'test')
datagen = ImageDataGenerator(rescale=1./255)
batch_size = 20
def extract_features(directory, sample_count):
    features = np.zeros(shape=(sample_count, 4, 4, 512))
    labels = np.zeros(shape=(sample_count))
    generator = datagen.flow_from_directory(
        directory,
        target_size=(150, 150),
        batch_size=batch_size,
        class_mode='binary')
    i=0
    for inputs_batch, labels_batch in generator:
        features[i * batch_size : (i + 1) * batch_size] = features_batch
        labels[i * batch_size : (i + 1) * batch_size] = labels_batch
        i += 1
        if i * batch_size >= sample_count:
            break
    return features, labels
```



```
In [41]:
```

```
train_features, train_labels = extract_features(train_dir, 2000)
validation_features, validation_labels = extract_features(validation_dir, 1000)
test_features, test_labels = extract_features(test_dir, 1000)
```

Found 2000 images belonging to 2 classes.
 Found 1000 images belonging to 2 classes.
 Found 1000 images belonging to 2 classes.


```
In [42]:
```

```
train_features = np.reshape(train_features, (2000, 4*4* 512))
validation_features = np.reshape(validation_features, (1000, 4*4* 512))
test_features = np.reshape(test_features, (1000, 4*4* 512))
```

Defining and training the densely connected classifier:

In [43]:

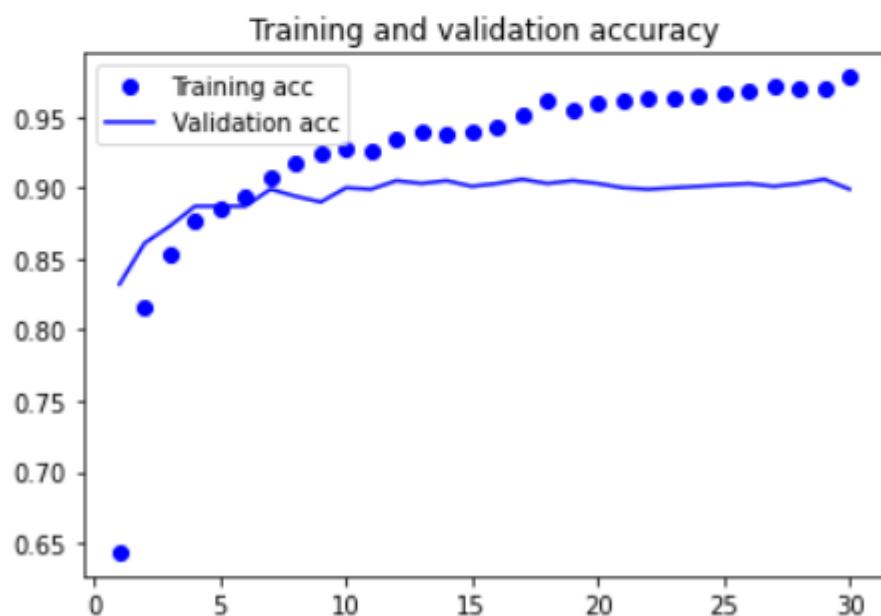
```
from keras import models
from keras import layers
from keras import optimizers
model = models.Sequential()
model.add(layers.Dense(256, activation='relu', input_dim=4 * 4 * 512))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(1, activation='sigmoid'))
model.compile(optimizer=optimizers.RMSprop(lr=2e-5),
loss='binary_crossentropy',
metrics=['acc'])
history = model.fit(train_features, train_labels,
epochs=30,
batch_size=20,
validation_data=(validation_features, validation_labels))
```

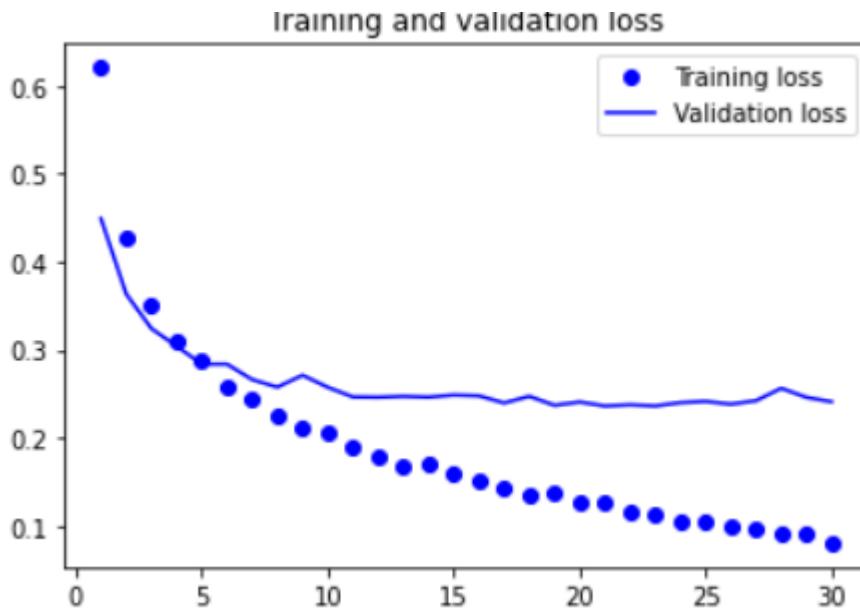
```
Epoch 1/30
100/100 [=====] - 4s 34ms/step - loss: 0.7039 - acc: 0.5675 - val_loss: 0.4490 - val_acc: 0.8320
Epoch 2/30
100/100 [=====] - 3s 32ms/step - loss: 0.4479 - acc: 0.8088 - val_loss: 0.3634 - val_acc: 0.8610
Epoch 3/30
100/100 [=====] - 3s 32ms/step - loss: 0.3646 - acc: 0.8445 - val_loss: 0.3242 - val_acc: 0.8730
Epoch 4/30
100/100 [=====] - 3s 32ms/step - loss: 0.3226 - acc: 0.8714 - val_loss: 0.3042 - val_acc: 0.8870
Epoch 5/30
100/100 [=====] - 3s 33ms/step - loss: 0.2883 - acc: 0.8913 - val_loss: 0.2837 - val_acc: 0.8870
Epoch 6/30
100/100 [=====] - 3s 32ms/step - loss: 0.2617 - acc: 0.8945 - val_loss: 0.2838 - val_acc: 0.8870
Epoch 7/30
100/100 [=====] - 3s 32ms/step - loss: 0.2373 - acc: 0.9174 - val_loss: 0.2662 - val_acc: 0.8990
Epoch 8/30
100/100 [=====] - 3s 33ms/step - loss: 0.2284 - acc: 0.9120 - val_loss: 0.2576 - val_acc: 0.8940
Epoch 9/30
100/100 [=====] - 3s 33ms/step - loss: 0.2102 - acc: 0.9211 - val_loss: 0.2712 - val_acc: 0.8900
Epoch 10/30
100/100 [=====] - 3s 33ms/step - loss: 0.2053 - acc: 0.9292 - val_loss: 0.2574 - val_acc: 0.9000
Epoch 11/30
100/100 [=====] - 3s 34ms/step - loss: 0.1834 - acc: 0.9359 - val_loss: 0.2465 - val_acc: 0.8990
Epoch 12/30
100/100 [=====] - 3s 30ms/step - loss: 0.1810 - acc: 0.9360 - val_loss: 0.2461 - val_acc: 0.9050
Epoch 13/30
100/100 [=====] - 3s 30ms/step - loss: 0.1689 - acc: 0.9437 - val_loss: 0.2471 - val_acc: 0.9030
Epoch 14/30
100/100 [=====] - 3s 30ms/step - loss: 0.1681 - acc: 0.9436 - val_loss: 0.2463 - val_acc: 0.9050
Epoch 15/30
100/100 [=====] - 3s 29ms/step - loss: 0.1572 - acc: 0.9384 - val_loss: 0.2486 - val_acc: 0.9010
```

Plotting the results:

In [44]:

```
import matplotlib.pyplot as plt
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```





Adding a densely connected classifier on top of the convolutional base:

```
In [45]: from keras import models
from keras import layers
model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

```
In [46]: model.summary()

Model: "sequential_3"
-----
```

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 4, 4, 512)	14714688
flatten_2 (Flatten)	(None, 8192)	0
dense_6 (Dense)	(None, 256)	2097408
dense_7 (Dense)	(None, 1)	257

```
-----
```

```
Total params: 16,812,353
Trainable params: 16,812,353
Non-trainable params: 0
```

```
In [47]: print('This is the number of trainable weights '
      'before freezing the conv base:', len(model.trainable_weights))
conv_base.trainable = False
print('This is the number of trainable weights '
      'after freezing the conv base:', len(model.trainable_weights))
```

This is the number of trainable weights before freezing the conv base: 30
 This is the number of trainable weights after freezing the conv base: 4

Training the model end to end with a frozen convolutional base:

In [48]:

```
from keras.preprocessing.image import ImageDataGenerator
from keras import optimizers
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest')
test_datagen = ImageDataGenerator(rescale=1./255)
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150),
    batch_size=20,
    class_mode='binary')
validation_generator = test_datagen.flow_from_directory(
    validation_dir,
    target_size=(150, 150),
    batch_size=20,
    class_mode='binary')
```

```
Found 2000 images belonging to 2 classes.
Found 1000 images belonging to 2 classes.
```

In [49]:

```
model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=2e-5),
              metrics=['acc'])
history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=30,
    validation_data=validation_generator,
    validation_steps=50)

Epoch 1/30
100/100 [=====] - 179s 2s/step - loss: 0.6106 - acc: 0.6838 - val_loss: 0.4186 - val_acc: 0.8330
Epoch 2/30
100/100 [=====] - 180s 2s/step - loss: 0.4734 - acc: 0.7778 - val_loss: 0.3546 - val_acc: 0.8600
Epoch 3/30
100/100 [=====] - 173s 2s/step - loss: 0.4129 - acc: 0.8182 - val_loss: 0.3174 - val_acc: 0.8760
Epoch 4/30
100/100 [=====] - 170s 2s/step - loss: 0.3877 - acc: 0.8387 - val_loss: 0.3060 - val_acc: 0.8700
Epoch 5/30
100/100 [=====] - 170s 2s/step - loss: 0.3893 - acc: 0.8366 - val_loss: 0.2889 - val_acc: 0.8800
Epoch 6/30
100/100 [=====] - 168s 2s/step - loss: 0.3760 - acc: 0.8335 - val_loss: 0.2842 - val_acc: 0.8800
Epoch 7/30
100/100 [=====] - 167s 2s/step - loss: 0.3495 - acc: 0.8540 - val_loss: 0.2758 - val_acc: 0.8810
Epoch 8/30
100/100 [=====] - 174s 2s/step - loss: 0.3282 - acc: 0.8600 - val_loss: 0.2706 - val_acc: 0.8890
Epoch 9/30
100/100 [=====] - 176s 2s/step - loss: 0.3339 - acc: 0.8643 - val_loss: 0.2584 - val_acc: 0.8960
Epoch 10/30
100/100 [=====] - 179s 2s/step - loss: 0.3385 - acc: 0.8568 - val_loss: 0.2573 - val_acc: 0.8990
Epoch 11/30
100/100 [=====] - 181s 2s/step - loss: 0.3218 - acc: 0.8642 - val_loss: 0.2540 - val_acc: 0.8980
Epoch 12/30
100/100 [=====] - 187s 2s/step - loss: 0.3277 - acc: 0.8583 - val_loss: 0.2655 - val_acc: 0.8860
Epoch 13/30
100/100 [=====] - 186s 2s/step - loss: 0.3349 - acc: 0.8570 - val_loss: 0.2507 - val_acc: 0.8960
Epoch 14/30
100/100 [=====] - 181s 2s/step - loss: 0.3153 - acc: 0.8660 - val_loss: 0.2522 - val_acc: 0.8970
Epoch 15/30
100/100 [=====] - 174s 2s/step - loss: 0.3028 - acc: 0.8720 - val_loss: 0.2474 - val_acc: 0.9060
```

```
In [50]: conv_base.summary()

Model: "vgg16"
_________________________________________________________________
Layer (type)                 Output Shape              Param #
=================================================================
input_1 (InputLayer)          [(None, 150, 150, 3)]   0
block1_conv1 (Conv2D)         (None, 150, 150, 64)    1792
block1_conv2 (Conv2D)         (None, 150, 150, 64)    36928
block1_pool (MaxPooling2D)    (None, 75, 75, 64)     0
block2_conv1 (Conv2D)         (None, 75, 75, 128)    73856
block2_conv2 (Conv2D)         (None, 75, 75, 128)    147584
block2_pool (MaxPooling2D)    (None, 37, 37, 128)    0
block3_conv1 (Conv2D)         (None, 37, 37, 256)    295168
block3_conv2 (Conv2D)         (None, 37, 37, 256)    590080
block3_conv3 (Conv2D)         (None, 37, 37, 256)    590080
block3_pool (MaxPooling2D)    (None, 18, 18, 256)    0
block4_conv1 (Conv2D)         (None, 18, 18, 512)    1180160
block4_conv2 (Conv2D)         (None, 18, 18, 512)    2359808
block4_conv3 (Conv2D)         (None, 18, 18, 512)    2359808
block4_pool (MaxPooling2D)    (None, 9, 9, 512)     0
block5_conv1 (Conv2D)         (None, 9, 9, 512)     2359808
block5_conv2 (Conv2D)         (None, 9, 9, 512)     2359808
block5_conv3 (Conv2D)         (None, 9, 9, 512)     2359808
block5_pool (MaxPooling2D)    (None, 4, 4, 512)     0
=================================================================
Total params: 14,714,688
Trainable params: 0
Non-trainable params: 14,714,688
```

Freezing all layers up to a specific one:

```
In [51]: conv_base.trainable = True
set_trainable = False
for layer in conv_base.layers:
    if layer.name == 'block5_conv1':
        set_trainable = True
    if set_trainable:
        layer.trainable = True
    else:
        layer.trainable = False

In [52]: model.compile(loss='binary_crossentropy',
                    optimizer=optimizers.RMSprop(lr=1e-5),
                    metrics=['acc'])
history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=100,
    validation_data=validation_generator,
    validation_steps=50)

Epoch 1/100
100/100 [=====] - 227s 2s/step - loss: 0.3005 - acc: 0.8740 - val_loss: 0.2392 - val_acc: 0.9130
Epoch 2/100
100/100 [=====] - 208s 2s/step - loss: 0.2596 - acc: 0.8911 - val_loss: 0.2114 - val_acc: 0.9170
Epoch 3/100
100/100 [=====] - 206s 2s/step - loss: 0.2603 - acc: 0.8926 - val_loss: 0.2096 - val_acc: 0.9120
Epoch 4/100
100/100 [=====] - 205s 2s/step - loss: 0.2133 - acc: 0.9082 - val_loss: 0.2364 - val_acc: 0.9050
Epoch 5/100
100/100 [=====] - 205s 2s/step - loss: 0.2028 - acc: 0.9126 - val_loss: 0.2022 - val_acc: 0.9240
Epoch 6/100
100/100 [=====] - 204s 2s/step - loss: 0.1966 - acc: 0.9227 - val_loss: 0.1837 - val_acc: 0.9280
Epoch 7/100
100/100 [=====] - 207s 2s/step - loss: 0.1729 - acc: 0.9307 - val_loss: 0.1961 - val_acc: 0.9300
Epoch 8/100
100/100 [=====] - 216s 2s/step - loss: 0.1623 - acc: 0.9353 - val_loss: 0.1883 - val_acc: 0.9300
Epoch 9/100
100/100 [=====] - 229s 2s/step - loss: 0.1672 - acc: 0.9360 - val_loss: 0.2084 - val_acc: 0.9250
Epoch 10/100
100/100 [=====] - 228s 2s/step - loss: 0.1507 - acc: 0.9469 - val_loss: 0.2007 - val_acc: 0.9330
```

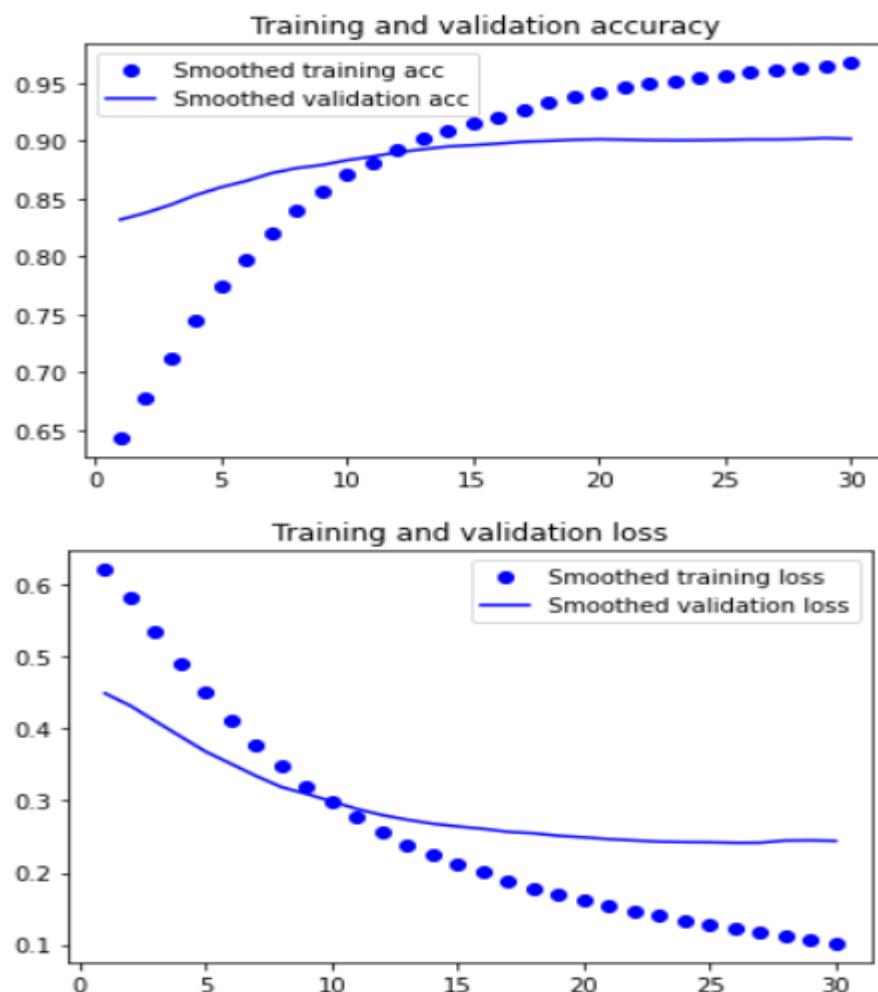
Smoothing the plots:

In [54]:

```
def smooth_curve(points, factor=0.8):
    smoothed_points = []
    for point in points:
        if smoothed_points:
            previous = smoothed_points[-1]
            smoothed_points.append(previous * factor + point * (1 - factor))
        else:
            smoothed_points.append(point)
    return smoothed_points

plt.plot(epochs,
smooth_curve(acc), 'bo', label='Smoothed training acc')
plt.plot(epochs,
smooth_curve(val_acc), 'b', label='Smoothed validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()

plt.plot(epochs,
smooth_curve(loss), 'bo', label='Smoothed training loss')
plt.plot(epochs,
smooth_curve(val_loss), 'b', label='Smoothed validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



Evaluate the model:

```
In [55]: test_generator = test_datagen.flow_from_directory(  
    test_dir,  
    target_size=(150, 150),  
    batch_size=20,  
    class_mode='binary')  
test_loss, test_acc = model.evaluate_generator(test_generator, steps=50)  
print('test acc:', test_acc)
```

Found 1000 images belonging to 2 classes.

test acc: 0.9190000295639038

Visualizing what convnets learn:

```
In [57]: from keras.models import load_model  
model = load_model('cats_and_dogs_small_2.h5')  
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_5 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_5 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_6 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_6 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_7 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_7 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dropout (Dropout)	(None, 6272)	0
dense_2 (Dense)	(None, 512)	3211776
dense_3 (Dense)	(None, 1)	513

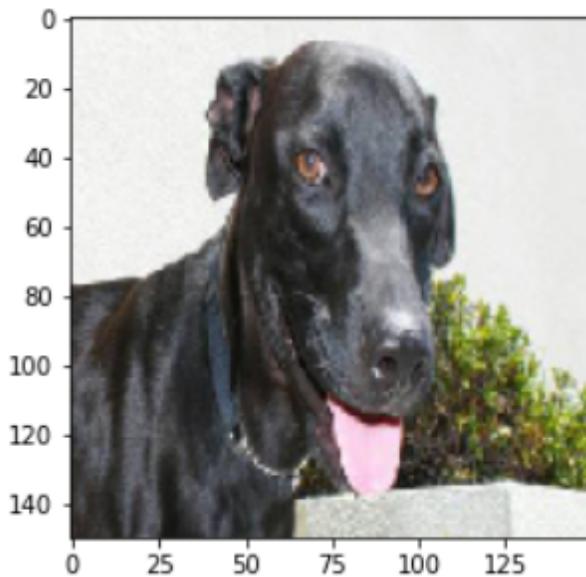
Total params: 3,453,121
Trainable params: 3,453,121
Non-trainable params: 0

Preprocessing a single image:

```
In [61]: img_path = 'C:/Users/ACER/Downloads/test1/23.jpg'  
from keras.preprocessing import image  
import numpy as np  
img = image.load_img(img_path, target_size=(150, 150))  
img_tensor = image.img_to_array(img)  
img_tensor = np.expand_dims(img_tensor, axis=0)  
img_tensor /= 255.  
  
print(img_tensor.shape)
```

(1, 150, 150, 3)

```
In [62]: import matplotlib.pyplot as plt  
plt.imshow(img_tensor[0])  
plt.show()
```



Instantiating a model from an input tensor and a list of output tensors and Visualizing:

```
In [63]: from keras import models
layer_outputs = [layer.output for layer in model.layers[:8]]
activation_model = models.Model(inputs=model.input, outputs=layer_outputs)
```

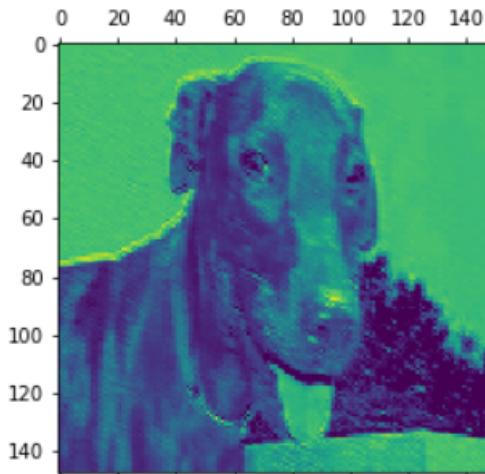
```
In [64]: activations = activation_model.predict(img_tensor)
```

```
In [65]: first_layer_activation = activations[0]
print(first_layer_activation.shape)
```

```
(1, 148, 148, 32)
```

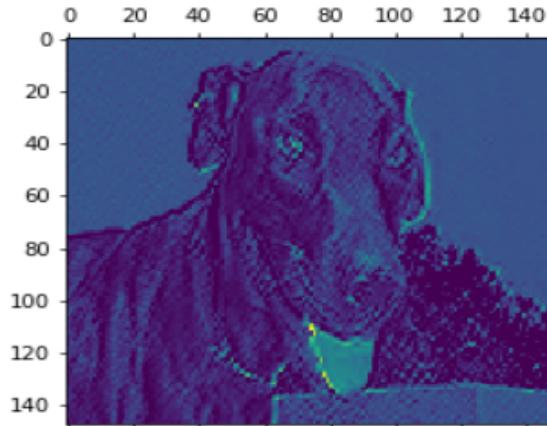
```
In [66]: import matplotlib.pyplot as plt
plt.matshow(first_layer_activation[0, :, :, 4], cmap='viridis')
```

```
Out[66]: <matplotlib.image.AxesImage at 0x18a02d28490>
```



```
In [67]: plt.matshow(first_layer_activation[0, :, :, 7], cmap='viridis')
```

```
Out[67]: <matplotlib.image.AxesImage at 0x18a02d09d30>
```



Visualizing convnet filters

Defining the loss tensor for filter visualization:

In [78]:

```
from keras.applications import VGG16
from keras import backend as K

model = VGG16(weights='imagenet',
include_top=False)
layer_name = 'block3_conv1'
filter_index = 0
layer_output = model.get_layer(layer_name).output
loss = K.mean(layer_output[:, :, :, :, filter_index])
```

In [79]:

```
grads = K.gradients(loss, model.input)[0]
```

In [80]:

```
grads /= (K.sqrt(K.mean(K.square(grads))) + 1e-5)
```

In [81]:

```
iterate = K.function([model.input], [loss, grads])
import numpy as np
loss_value, grads_value = iterate([np.zeros((1, 150, 150, 3))])
```

In [83]:

```
input_img_data = np.random.random((1, 150, 150, 3)) * 20 + 128.
step = 1.
for i in range(40):
    loss_value, grads_value = iterate([input_img_data])
    input_img_data += grads_value * step
```

In [84]:

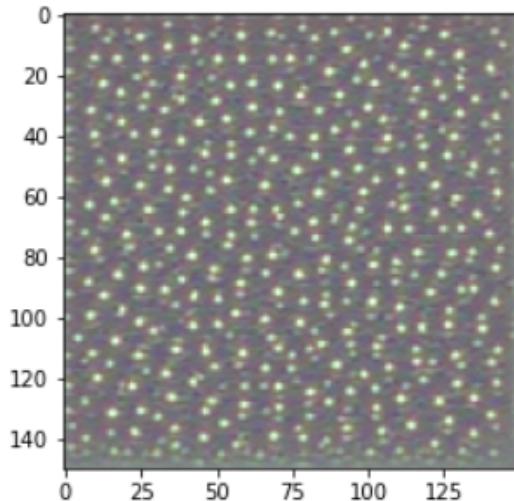
```
def deprocess_image(x):
    x -= x.mean()
    x /= (x.std() + 1e-5)
    x *= 0.1
    x += 0.5
    x = np.clip(x, 0, 1)
    x *= 255
    x = np.clip(x, 0, 255).astype('uint8')
    return x
```

Function to generate filter visualizations

```
In [86]: def generate_pattern(layer_name, filter_index, size=150):
    layer_output = model.get_layer(layer_name).output
    loss = K.mean(layer_output[:, :, :, filter_index])
    grads = K.gradients(loss, model.input)[0]
    grads /= (K.sqrt(K.mean(K.square(grads))) + 1e-5)
    iterate = K.function([model.input], [loss, grads])
    input_img_data = np.random.random((1, size, size, 3)) * 20 + 128.
    step = 1.
    for i in range(40):
        loss_value, grads_value = iterate([input_img_data])
        input_img_data += grads_value * step
        img = input_img_data[0]
    return deprocess_image(img)
```

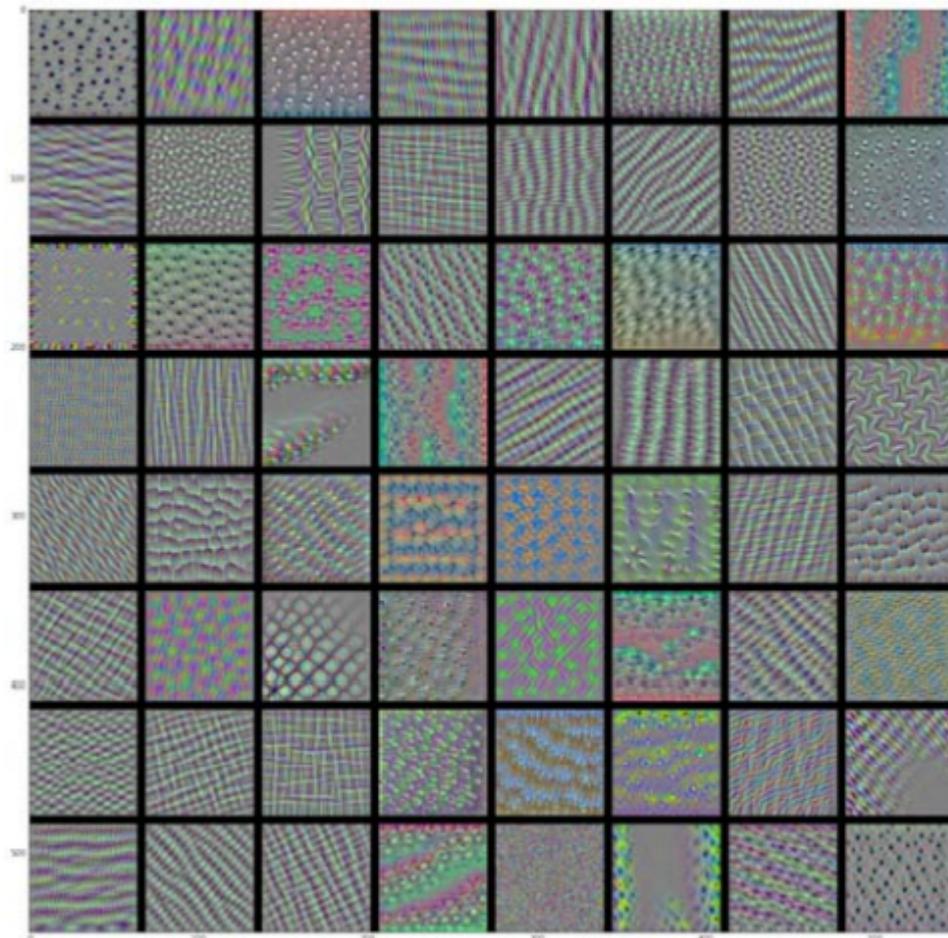
```
In [87]: plt.imshow(generate_pattern('block3_conv1', 0))
```

Out[87]: <matplotlib.image.AxesImage at 0x18a075113a0>



Generating a grid of all filter response patterns in a layer:

```
In [89]: layer_name = 'block1_conv1'
size = 64
margin = 5
results = np.zeros((8 * size+7* margin, 8 * size+7* margin, 3))
for i in range(8):
    for j in range(8):
        filter_img = generate_pattern(layer_name, i + (j * 8), size=size)
        horizontal_start = i * size + i * margin
        horizontal_end = horizontal_start + size
        vertical_start = j * size + j * margin
        vertical_end = vertical_start + size
        results[horizontal_start: horizontal_end,
               vertical_start: vertical_end, :] = filter_img
plt.figure(figsize=(20, 20))
plt.imshow(results)
```



Loading the VGG16 network with pretrained weights:

```
In [90]:  
from keras.applications.vgg16 import VGG16  
model = VGG16(weights='imagenet')  
  
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels.h5  
553467904/553467096 [=====] - 2104s 4us/step
```

```
In [91]:  
from keras.preprocessing import image  
from keras.applications.vgg16 import preprocess_input, decode_predictions  
import numpy as np  
img_path = 'C:/Users/ACER/Downloads/test1/50.jpg'  
img = image.load_img(img_path, target_size=(224, 224))  
x = image.img_to_array(img)  
x = np.expand_dims(x, axis=0)  
x = preprocess_input(x)
```

```
In [92]:  
preds = model.predict(x)  
print('Predicted:', decode_predictions(preds, top=3)[0])
```

```
Downloading data from https://storage.googleapis.com/download.tensorflow.org/data/imagenet_class_index.json
40960/35363 [=====] - 0s 3us/step
Predicted: [('n02105855', 'Shetland_sheepdog', 0.167318), ('n02110185', 'Siberian_husky', 0.11742018), ('n02106166', 'Border_collie', 0.039714873)]
```

```
In [93]: np.argmax(preds[0])
```

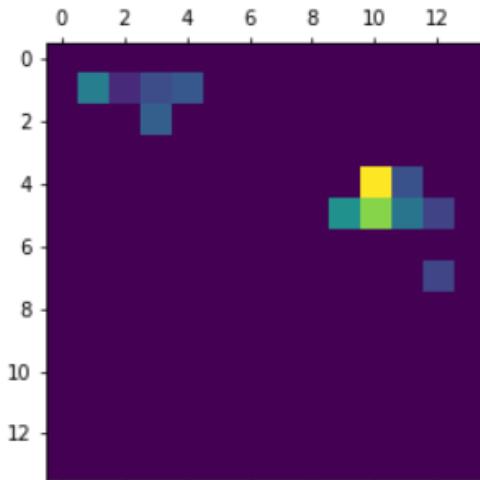
```
Out[93]: 230
```

Setting up the Grad-CAM algorithm:

```
In [95]: cat = model.output[:, 386]
last_conv_layer = model.get_layer('block5_conv3')
grads = K.gradients(cat, last_conv_layer.output)[0]
pooled_grads = K.mean(grads, axis=(0, 1, 2))
iterate = K.function([model.input],
[pooled_grads, last_conv_layer.output[0]])
pooled_grads_value, conv_layer_output_value = iterate([x])
for i in range(512):
    conv_layer_output_value[:, :, i] *= pooled_grads_value[i]
    heatmap = np.mean(conv_layer_output_value, axis=-1)
```

```
In [96]: heatmap = np.maximum(heatmap, 0)
heatmap /= np.max(heatmap)
plt.matshow(heatmap)
```

```
Out[96]: <matplotlib.image.AxesImage at 0x18a0a808040>
```



```
In [100... import cv2
img = cv2.imread(img_path)
heatmap = cv2.resize(heatmap, (img.shape[1], img.shape[0]))
heatmap = np.uint8(255 * heatmap)
heatmap = cv2.applyColorMap(heatmap, cv2.COLORMAP_JET)
superimposed_img = heatmap * 0.4 + img
cv2.imwrite('C:/Users/ACER/Downloads/test1/50.jpg', superimposed_img)
```

```
Out[100... True
```



Results:

- From the accuracy graphs VGG16(96%) pre-trained model performed well compared to the convnet(82%) that is trained.

Conclusion:

Thus cat and dog images are classified using convnets and VGG16 models and accuracy is evaluated. Data Augmentation was used to avoid overfitting. Filters learned by the convnets are visualized.

References:

- [How to Classify Photos of Dogs and Cats \(with 97% accuracy\) \(machinelearningmastery.com\)](#)
- [Cats And Dogs Image Classification Using Keras – Pythonista Planet](#)
- [Keras Implementation of VGG16 Architecture from Scratch with Dogs Vs Cat Data Set | MLK - Machine Learning Knowledge](#)